

Lattice QCD on Intel Xeon Phi's

Peter Labus

in collaboration with Bartosz Kostrzewa and Martin Ueding
@ HISKP Universität Bonn

JSC Jülich
10.02.2016

LQCD & the rôle of the *dslash* stencil

- ▶ Calculate integrals of the form

$$\int \mathcal{D}\Phi f[\Phi] e^{-S_{QCD}[\Phi]}$$

after *discretising* space-time through a lattice

- ▶ Algorithmic layers of LQCD:
 - ❶ Hybrid Monte Carlo
 - ❷ Efficient Krylov Solvers for $Mx = b$ in Molecular Dynamics
 - ❸ BLAS linear algebra
- ▶ Most *expensive* part:
 - Matrix times Vector (very large 10^9 squared & very sparse), described by (next-to) nearest-neighbour stencil

The QPhiX library (mainly: Balint Joo, Jefferson Lab)

- ▶ Aim: provide stencil operations for **general vector machines**
- ▶ Parallelised via QMP + OpenMP + SIMD (vector intrinsics)
- ▶ C++11 template library with external C++ code generator
(~ 25k lines of code + 10k testing & timing)
- ▶ Implements
 - ❶ Wilson / Wilson-Clover *dslash* stencils
 - ❷ BLAS linear algebra
 - ❸ Krylov Solvers (CG, BiCGStab, Mixed Precision, MultiCG)
- ▶ Common template parameters:
 - ❶ `typename FT`
 - ❷ `int VECLLEN`
 - ❸ `int SOALEN`
 - ❹ `bool COMPRESS12`
- ▶ **Our goal: Integrate Twisted-Mass Fermions in QPhiX**
This amounts to a change in the Fermion Matrix $M \dots$

The Fermion Matrix

Full (even-odd preconditioned) matrix is

$$M = A - \frac{1}{4} \not{D} A^{-1} \not{D},$$

where \not{D} is the nearest-neighbour ***ds*lash stencil**.

M can be constructed from two BLAS Matrix-Vector *kernels*:

$$(1) \ A^{-1} \not{D} x, \quad (2) \ A x - \frac{1}{4} \not{D} y.$$

The vectors x and y are called *spinors*.

The **Clover Term** A includes the **new** *twisted-mass*:

$$A \equiv \alpha \mathbb{1} + c_{sw} T \pm i\mu\gamma_5,$$

where c_{sw} and μ are user parameters, while α is fixed.

The *dslash* stencil in action

$$\begin{aligned}\chi_{\alpha}^a(x) &= \sum_{y \in \Lambda} \sum_{b=0}^2 \sum_{\beta=0}^3 \mathcal{D}_{\alpha\beta}^{ab}(x, y) \psi_{\beta}^b(y) \\ &= \sum_{b=0}^2 \sum_{\beta=0}^3 \sum_{\mu=0}^3 \left[U^{ab}(x, x + \hat{\mu}) (\mathbb{1} - \gamma_{\mu})_{\alpha\beta} \psi_{\beta}^b(x + \hat{\mu}) + \right. \\ &\quad \left. U^{\dagger ab}(x, x - \hat{\mu}) (\mathbb{1} + \gamma_{\mu})_{\alpha\beta} \psi_{\beta}^b(x - \hat{\mu}) \right]\end{aligned}$$

- ▶ x, y points/sites of the lattice Λ
- ▶ ψ input **spinor**: $3 \times 4 \times 2 = 24$ components per site
- ▶ χ output spinor
- ▶ U **gauge** or link variables: $3 \times 3 \times 2 = 18$ components per site
- ▶ γ_{μ} : 4 constant 4×4 complex matrices
- ▶ $\hat{\mu}$: 4 unit vectors (one for each space-time dimension)

The Clover Term

A is a 12×12 complex, block-diagonal matrix (each lattice site):

$$A = \begin{pmatrix} (\alpha \pm i\mu) \mathbb{1}_6 + c_{SW} B_0 & \mathbb{O}_6 \\ \mathbb{O}_6 & (\alpha \mp i\mu) \mathbb{1}_6 + c_{SW} B_1 \end{pmatrix}$$

QPhiX so far:

- ▶ $\mu = 0 \Rightarrow A$ and A^{-1} are hermitian (and block-diagonal)
- ▶ need to store only two blocks with
 - ▶ 6 real diagonal elements
 - ▶ 15 complex off-diagonal elements

We want: $\mu \neq 0$

- ▶ then A is not hermitian
- ▶ in particular A^{-1} is dense (but still block-diagonal)
- ▶ need to store two full 6×6 complex blocks

Arithmetic Intensities

$$I_A = \text{FLOPs per Byte of moved data}$$

dslash stencil: $\nabla^2 x$

1320 FLOPs per site

$$I_A = 1.06 \text{ FLOP/Byte}$$

combined multiplication with clover term: $A^{-1} \nabla^2 x$

1872 FLOPs per site

$$\mu = 0 : I_A = 1.17 \text{ FLOP/Byte}$$

$$\mu \neq 0 : I_A = 0.98 \text{ FLOP/Byte}$$


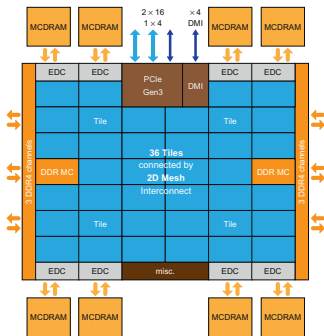
Theoretical Peaks on KNL / KNC:

$$I_A = 15.4 / 6.3 \text{ FLOP/Byte}$$

ALL ROUTINES ARE MEMORY BANDWIDTH BOUND
Expect about 20% less performance with twisted-mass

Combining these Intensities with a simple hardware model and the (measured) memory bandwidth gives us a good *rooftop model* ...

Intel Xeon Phi Knights Landing (KNL)



- ▶ up to 72 cores @ 1.5 GHz (3.5 TFLOP/s)
- ▶ 2D Mesh: 700 GB/s
- ▶ 8 GB MCDRAM: 450 GB/s
- ▶ 384 GB DDR-4: 90 GB/s

- ▶ 2 cores / 4 VPUs
- ▶ 512-bit Vector Registers (AVX-512)
- ▶ 4 hardware hyper-threads
- ▶ 1 MB shared L2 cache

General Programming Implications

Three Guidelines...

- 1 **Scaling**
- 2 **Vectorisation**
- 3 **Data Locality & Cache Re-use**

... and how they are implemented in QPhiX

- 1 **OpenMP thread scheduling** for lattice traversal
- 2 Multi-ISA **C++ code generator** using Vector Intrinsics
- 3 **Cache blocking** into L2 and **Structures of Arrays** data layout (*tiles*)


Optimisations in QPhiX-codegen:

L1 & L2 software prefetches, dimensions of tiles (SOALEN), streaming stores, gather/scatter, ...

Data layout & Cache blocking

- ▶ Elementary vector elements are *Structures of Arrays* (SoA):

```
typedef FT FourSpinorBlock[3][4][2][SOALEN];
```
- ▶ SOALEN may be a factor of VECLEN = $\text{ngy} * \text{SOALEN}$:
form *tiles* in X-Y-plane (Xeon Phi SP: 16×1 , 8×2 , 4×4)
- ▶ Then divide lattice in blocks:
X: SOALEN
Y, Z: B_y , B_z
such that 3 T-Slices fit into L2
- ▶ Distribute blocks to
cores/threads & process in
chunks of SIMD vectors, *i.e.* tiles



QPhiX-codegen

- ▶ Three abstract (class) objects:
 - 1 Instruction: FLOPs, Memory inst., scope delimiters, if-else, ...
 - 2 Address: Memory addresses to deal with pointers & offsets
 - 3 FVec: “The registers” which are fed to intrinsic calls, *i.e.* variables in C++ code (*e.g.* of type `_mm512d`)
- ▶ First two classes: `std::string serialize(void)`
- ▶ Save Instructions into `InstVector`’s and dump to file

```
for( auto ops : kernels ) {  
  
    InstVector ivector;  
    InstVector l2prefs;  
    std::ostringstream filename;  
  
    filename << ...;  
  
    // Generate Instructions  
    generateL2Prefetches(l2prefs, compress12, chi-prefetches, clover, twisted_mass);  
    dslash_plain_body(ivector, compress12, clover, twisted_mass, isPlus);  
    mergeIvectorWithL2Prefetches(ivector, l2prefs);  
    dumpIvector(ivector, filename.str());  
  
}
```

Code Example 1: NEW Clover Term Multiplication

```
typedef struct {
    FT diag1[6][VECLEN];
    FT off_diag1[15][2][VECLEN];
    FT diag2[6][VECLEN];
    FT off_diag2[15][2][VECLEN];
} Clover;

typedef struct {
    FT block1[6][6][2][VECLEN];
    FT block2[6][6][2][VECLEN];
} FullClover;

void full_clover_term(InstVector& ivector, FVec in_spinor[4][3][2], bool face, std::string mask)
{
    for(int block=0; block!=2; ++block) {

        PrefetchL1FullCloverFullBlockIn(ivector, clBase, clOffs, block);
        LoadFullCloverFullBlock(ivector, clov_full, clBase, clOffs, block);

        for(int sc1=0; sc1!=6; ++sc1) { // half-spin-colour row

            // Calculate out indices
            FVec *clout = out_spinor[spin_out][col_out];

            for(int sc2=0; sc2!=6; ++sc2) { // half-spin-colour column

                // Calculate in indices
                FVec *clin = in_spinor[spin_in][col_in];

                if(sc2 == 0 && !face)    mulCVec(ivector, clout, clov_full[sc1][sc2], clin, mask);
                else                    fmaddCVec(ivector, clout, clov_full[sc1][sc2], clin, clout, mask);
            }
        }
    }
}
```

Code Example 2: Extending the User Interface

```
template <typename FT, int veclen, int soalen, bool compress12>
class TMClovDslash {
public:
    // Constructor (padded and aligned data allocation)
    TMClovDslash(Geometry<FT, veclen, soalen, compress12>* geom_,
        double t_boundary_,
        double dslash_aniso_s_,
        double dslash_aniso_t_);

    // Destructor
    ~TMClovDslash();


    // Geometry Getter
    Geometry<FT, veclen, soalen, compress12>& getGeometry(void) { return (*s); }

    //  $A^{-1} * D * \psi$ 
    // (lattice traversal w/ pointer and offset calculation & MPI comm)
    void dslash(FourSpinorBlock* res,
        const FourSpinorBlock* psi,
        const SU3MatrixBlock* u,
        const FullCloverBlock* invclov[2],
        int isign,
        int cb);


    //  $A * \chi - b * D * \psi$ 
    void dslashAChiMinusBDPsi(FourSpinorBlock* res,
        const FourSpinorBlock* psi,
        const FourSpinorBlock* chi,
        const SU3MatrixBlock* u,
        const FullCloverBlock* clov[2],
        const double beta,
        int isign,
        int cb);
}
```

Performance of $A^{-1} \nabla x$ Kernels


$c_{sw} = 0, \mu = 0$




$c_{sw} = 0, \mu \neq 0$



$c_{sw} \neq 0, \mu = 0$




$c_{sw} \neq 0, \mu \neq 0$




Optimisation Features on KNC vs. KNL

Displayed is the basic *dslash* stencil ($c_{sw} = \mu = 0$) in single precision with $SOA = 16$.

KNC:




KNL:




- ▶ KNL natively supports prefetches & streaming stores much better than KNC
- ▶ Visible benefits from data layout tuning (shown here: SoA length tuning)

Multi-node results

Dual-Socket Xeon (Haswell):



Xeon Phi (KNC):



- ▶ STRONG scaling up to 64 Xeon / 16 KNC nodes for various volumes
- ▶ WEAK scaling up to 64 KNC nodes

Conclusions & Outlook

- ▶ Added new data type & dynamic memory allocation facilities
- ▶ Implemented low-level Twisted-Mass-Clover kernels & hardware feature utility functions
- ▶ Extended User Interface, added Testing & Timing for new kernels, expanded autoconf/automake setup ...
- ▶ Confirmed performance expectations on single KNL's processors

Conclusions & Outlook

- ▶ Added new data type & dynamic memory allocation facilities
- ▶ Implemented low-level Twisted-Mass-Clover kernels & hardware feature utility functions
- ▶ Extended User Interface, added Testing & Timing for new kernels, expanded autoconf/automake setup ...
- ▶ Confirmed performance expectations on single KNL's processors

- ▶ Interface for tmLQCD
- ▶ Non-degenerate Twisted-Mass/Twisted-Mass-Clover *dslash*
- ▶ Twisted boundary conditions
- ▶ New Algorithms: Domain decomposition, Multi-grid, ...?

Conclusions & Outlook

- ▶ Added new data type & dynamic memory allocation facilities
- ▶ Implemented low-level Twisted-Mass-Clover kernels & hardware feature utility functions
- ▶ Extended User Interface, added Testing & Timing for new kernels, expanded autoconf/automake setup ...
- ▶ Confirmed performance expectations on single KNL's processors

- ▶ Interface for tmLQCD
- ▶ Non-degenerate Twisted-Mass/Twisted-Mass-Clover *dslash*
- ▶ Twisted boundary conditions
- ▶ New Algorithms: Domain decomposition, Multi-grid, ...?

THANKS!